

Python socket receive - incoming packets always have a different size

 stackoverflow.com/questions/1708835/python-socket-receive-incoming-packets-always-have-a-different-size

37



I'm using the SocketServer module for a TCP server. I'm experiencing some issue here with the `recv()` function, because the incoming packets always have a different size, so if I specify `recv(1024)` (I tried with a bigger value, and smaller), it gets stuck after 2 or 3 requests because the packet length will be smaller (I think), and then the server gets stuck until a timeout.

```
class Test(SocketServer.BaseRequestHandler):

    def handle(self):

        print "From:", self.client_address

        while True:

            data = self.request.recv(1024)
            if not data: break

            if data[4] == "\x20":
                self.request.sendall("hello")
            if data[4] == "\x21":
                self.request.sendall("bye")
            else:
                print "unknow packet"
            self.request.close()
            print "Disconnected", self.client_address

launch = SocketServer.ThreadingTCPServer(('', int(sys.argv[1])),Test)

launch.allow_reuse_address= True;

launch.serve_forever()
```

If the client sends multiples requests over the same source port, but the server gets stuck, any help would be very appreciated, thank !

asked Nov 10 '09 at 15:34



n00bie

51322 gold badges99 silver badges1010 bronze badges

45



The network is *always* unpredictable. TCP makes a lot of this random behavior go away for you. One wonderful thing TCP does: it guarantees that the bytes will arrive in the same order. But! It does *not* guarantee that they will arrive chopped up in the same way. You simply *cannot* assume that every `send()` from one end of the connection will result in exactly one `recv()` on the far end with exactly the same number of bytes.

When you say `socket.recv(x)`, you're saying 'don't return until you've read x bytes from the socket'. This is called "blocking I/O": you will block (wait) until your request has been filled. If every message in your protocol was exactly 1024 bytes, calling `socket.recv(1024)` would work great. But it sounds like that's not true. If your messages are a fixed number of bytes, just pass that number in to `socket.recv()` and you're done.

But what if your messages can be of different lengths? The first thing you need to do: stop calling `socket.recv()` with an explicit number. Changing this:

```
data = self.request.recv(1024)
```

to this:

```
data = self.request.recv()
```

means `recv()` will always return whenever it gets new data.

But now you have a new problem: how do you know when the sender has sent you a complete message? The answer is: you don't. You're going to have to make the length of the message an explicit part of your protocol. Here's the best way: prefix every message with a length, either as a fixed-size integer (converted to network byte order using `socket.ntohs()` or `socket.ntohl()` please!) or as a string followed by some delimiter (like '123:'). This second approach often less efficient, but it's easier in Python.

Once you've added that to your protocol, you need to change your code to handle `recv()` returning arbitrary amounts of data at any time. Here's an example of how to do this. I tried writing it as pseudo-code, or with comments to tell you what to do, but it wasn't very clear. So I've written it explicitly using the length prefix as a string of digits terminated by a colon. Here you go:

```

length = None
buffer = ""
while True:
    data += self.request.recv()
    if not data:
        break
    buffer += data
    while True:
        if length is None:
            if ':' not in buffer:
                break
            # remove the length bytes from the front of buffer
            # leave any remaining bytes in the buffer!
            length_str, ignored, buffer = buffer.partition(':')
            length = int(length_str)

        if len(buffer) < length:
            break
        # split off the full message from the remaining bytes
        # leave any remaining bytes in the buffer!
        message = buffer[:length]
        buffer = buffer[length:]
        length = None
    # PROCESS MESSAGE HERE

```

answered Nov 11 '09 at 16:02



Larry Hastings

2,226 1616 silver badges 1111 bronze badges

130



The answer by Larry Hastings has some great general advice about sockets, but there are a couple of mistakes as it pertains to how the `recv(bufsize)` method works in the Python socket module.

So, to clarify, since this may be confusing to others looking to this for help:

1. The `bufsize` param for the `recv(bufsize)` method is not optional. You'll get an error if you call `recv()` (without the param).
2. The `bufferlen` in `recv(bufsize)` is a *maximum* size. The `recv` will happily return fewer bytes if there are fewer available.

See [the documentation](#) for details.

Now, if you're receiving data from a client and want to know when you've received all of the data, you're probably going to have to add it to your protocol -- as Larry suggests. See [this recipe](#) for strategies for determining end of message.

As that recipe points out, for some protocols, the client will simply disconnect when it's done sending data. In those cases, your `while True` loop should work fine. If the client does *not* disconnect, you'll need to figure out some way to signal your content length, delimit your messages, or implement a timeout.

I'd be happy to try to help further if you could post your exact client code and a description of your test protocol.

answered Nov 27 '09 at 5:43



Hans L

4,17444 gold badges1818 silver badges2020 bronze badges

17



You can alternatively use `recv(x_bytes, socket.MSG_WAITALL)`, which seems to work only on Unix, and will return exactly `x_bytes`.

answered Dec 2 '09 at 4:40



henrietta

17122 bronze badges

2



That's the nature of TCP: the protocol fills up packets (lower layer being IP packets) and sends them. You can have some degree of control over the MTU (Maximum Transfer Unit).

In other words: you must devise a protocol that rides on top of TCP where your "payload delineation" is defined. By "payload delineation" I mean the way you extract the unit of message your protocol supports. This can be as simple as "every NULL terminated strings".

answered Nov 10 '09 at 15:38



jldupont

74.1k4646 gold badges183183 silver badges298298 bronze badges

1



Note that **exact reason** why your code is frozen is **not** because you set too high request.recv() buffer size. Here is explained [What means buffer size in socket.recv\(buffer_size\)](#)

This code will work until it'll receive an **empty TCP message** (if you'd print this empty message, it'd show `b''`):

```
while True:
    data = self.request.recv(1024)
    if not data: break
```

And note, that there is **no way** to send empty TCP message. `socket.send(b'')` simply won't work.

Why? Because empty message is sent only when you type `socket.close()`, so your script will loop as long as you won't close your connection. As **Hans L** pointed out here are some [good methods to end message](#).

[edited Dec 15 '17 at 11:46](#)

[answered Dec 15 '17 at 11:36](#)



[Qback](#)

1,80511 gold badge99 silver badges2727 bronze badges

0



I know this is old, but I hope this helps someone.

Using regular python sockets I found that you can send and receive information in packets using sendto and recvfrom

```
# tcp_echo_server.py
import socket
```

```
ADDRESS = ''
PORT = 54321
```

```
connections = []
host = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
host.setblocking(0)
host.bind((ADDRESS, PORT))
host.listen(10) # 10 is how many clients it accepts
```

```
def close_socket(connection):
    try:
        connection.shutdown(socket.SHUT_RDWR)
    except:
        pass
    try:
```

```

        connection.close()
    except:
        pass

def read():
    for i in reversed(range(len(connections))):
        try:
            data, sender = connections[i][0].recvfrom(1500)
            return data
        except (BlockingIOError, socket.timeout, OSError):
            pass
        except (ConnectionResetError, ConnectionAbortedError):
            close_socket(connections[i][0])
            connections.pop(i)
    return b'' # return empty if no data found

def write(data):
    for i in reversed(range(len(connections))):
        try:
            connections[i][0].sendto(data, connections[i][1])
        except (BlockingIOError, socket.timeout, OSError):
            pass
        except (ConnectionResetError, ConnectionAbortedError):
            close_socket(connections[i][0])
            connections.pop(i)

# Run the main loop
while True:
    try:
        con, addr = host.accept()
        connections.append((con, addr))
    except BlockingIOError:
        pass

    data = read()
    if data != b'':
        print(data)
        write(b'ECHO: ' + data)
        if data == b'exit':
            break

# Close the sockets
for i in reversed(range(len(connections))):
    close_socket(connections[i][0])
    connections.pop(i)
close_socket(host)

```

The client is similar

```

# tcp_client.py
import socket

ADDRESS = "localhost"
PORT = 54321

s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.connect((ADDRESS, PORT))
s.setblocking(0)

def close_socket(connection):
    try:
        connection.shutdown(socket.SHUT_RDWR)
    except:
        pass
    try:
        connection.close()
    except:
        pass

def read():
    """Read data and return the read bytes."""
    try:
        data, sender = s.recvfrom(1500)
        return data
    except (BlockingIOError, socket.timeout, AttributeError, OSError):
        return b''
    except (ConnectionResetError, ConnectionAbortedError, AttributeError):
        close_socket(s)
        return b''

def write(data):
    try:
        s.sendto(data, (ADDRESS, PORT))
    except (ConnectionResetError, ConnectionAbortedError):
        close_socket(s)

while True:
    msg = input("Enter a message: ")
    write(msg.encode('utf-8'))

    data = read()
    if data != b'':
        print("Message Received:", data)

    if msg == "exit":
        break

close_socket(s)

```

answered Apr 26 '17 at 18:30



justengel

4,42522 gold badges1616 silver badges3535 bronze badges